Table of contents

This [repository](#) contains the supplementary materials, data processing pipelines, and scripts used for **'spanishoddata: A package for accessing and working with large Spanish mobility datasets'** article. The snapshot of this repository is at Zenodo: https://doi.org/10.5281/zenodo.15207374.

To open this repository online using interactive `RStudio` environment and run the scripts, just click the link >> https://mybinder.org/v2/gh/e-kotov/spanishoddata-paper-supplement/HEAD?urlpath=rstudio., then read the [relevant section below](#).

1 Getting Started

If you are only looking for the aritcle plots and supplement data (such as articles analysis data, search queries and results), kindly see the ["Repository Structure" -> "General interest files"](#) section. If you would like to learn more about the data processing pipelines, please refer to the ["Repository Structure" -> "Files for reproducibility"](#) section and the following sections that describe each pipeline.

2 Repository Structure

2.1 General interest files

Top-level folder / file	Second-level item	Description
plots/	main-plots/	Figures for the main article generated by the **"main"** pipeline
	supplement-plots/	Figures for the supplement generated by the **"main"** pipeline
	pipeline-plots/	Figures for the all `targets` pipelines used in the project.
supplement-data/	articles-using-mitms-mobility-data.json	Raw collected data with attributes extracted from the analysed articles
	articles-using-mitms-mobility-data.csv	Flattened table derived from the *.json* file
	fig_package_workflow.gv	`Graphviz` DOT file for the workflow diagram in the supplement
	scopus-search-query.sql	SQL query used to search Scopus
	scopus-search-results.bib	Raw Scopus search results in `BibTeX` format
	speed_test_summary.csv	Raw speed-test data produced by the **"speed_test"** pipeline (used in supplement)
qmd/	—	Human-readable Quarto/Markdown files to reproduce the case studies

2.2 Files for reproducibility

containers/	Dockerfile	`Dockerfile` to reproduce the computational environment
	update-mermaid-js.R	R script that updates `Mermaid.js` library that i used for workflow plot generation
R/	Subfolders: cache_osm_data, data_prep, main, speed_test, pipeline_plots	R scripts and supporting functions used across data-processing pipelines in respective targets pipelines
renv/	—	Folder holding the `renv` activation script and settings for in-

stalling the packages required to reproduce all code in the repository. To learn more about `renv` see https://rstudio.github.io/renv/.

renv.lock	—	Lockfile pinning exact R-package versions for `renv`.
.Rprofile	—	R startup file that enables `renv`.
_targets.yaml	—	Global configuration for the `targets` pipelines used in the project. To learn more about `targets` see https://books.ropensci.org/targets/ and https://docs.ropensci.org/targets/.
_targets_cache_osm_-data.R	—	`targets` script that caches OpenStreetMap data to make a snapshot. The cache is also stored at https://doi.org/10.5281/zenodo.15207222
_targets_data_prep.R	—	`targets` script for data-preparation steps. It downloads all data from original data sources and performs all of the preprocessing of data before passing it on to the "main" pipeline.
_targets_speed_test.R	—	`targets` script that benchmarks data processing speed between CSV, DuckDB and parquet formats and produces the **supplement-data/speed_test_summary.csv** file
_targets_main.R	—	`targets` script that creates all figures for the article.
_targets_pipeline_-plots.R	—	`targets` script that builds the figures for this repository that provide an overview of each pipeline.

Dockerfile	—	Dockerfile for https://mybinder.org/ or **docker2repo** to run the computational environment for reproducing all targets pipelines.

`Docker` and `Apptainer` container images with presereved computaitonal environment are deposited with a copy of this repository at https://doi.org/10.5281/zenodo.15207374.

3 Reproducing the analysis

3.1.1 Overview

This repository contains all files required to reproduce the analysis described in the article. Some analysis steps are inherently unreproducible, such as "cache_osm_data", as they take a snapshot of latest data at the time of running the code and later the retrieved data may change. Some steps require very large data downloads and processing power ("speed_test" and "data_prep"). Therefore these steps were isolated in separate pipelines. The repository comes with snapshots from these pipelines that are sufficient to run the "main" pipeline that reproduces all figures for the main article and the supplement, however if you would like to reproduce all other pipelines, be prepared for large data downloads and long processing times.

The "main" step that generates figures can be reproduced in a cloud hosted environment using the pre-build container image.

3.1.2 Option 1: Reproducing the figures in a web-browser using Binder
- Just click the link >> https://mybinder.org/v2/gh/e-kotov/spanishoddata-paper-supplement/HEAD?urlpath=rstudio.
- Wait for RStudio to load inside the web browser tab.
- Run `targets::tar_destroy` in R console to delete the pipeline snapshots.
- Run `targets::tar_make()` in R console to regenerate all figures.
- You will find the updated figures in the `plots/` folder.

You will likely not be able to run any other pipeline in this enviromnet, as it does not have enough compute power.

3.1.3 Option 2: Reproducing with locally installed `R`
Before proceeding with the steps make sure the computational environment is set up correctly.

- Install R 4.3.2 or later.

- Clone the repository or download it manually and unpack.

```
git clone https://github.com/e-kotov/spanishoddata-paper-supplement.git
```

- Run R or start `RStudio`/`Positron`/`VSCode` in the root directory of the project.

- Install packages with:

```
renv::restore(prompt = FALSE)
```

You can now proceed to the [pipelines section](#) below.

3.1.4 Option 3: `Docker` or `Apptainer` container

`Docker` and `Apptainer` Container images with complete computational environment to reproduce this project can be downloaded from https://doi.org/10.5281/zenodo.15207374. You also need to have [Docker](#) or `Docker`-compatible software installed to run the containers. For `Apptainer` container you can use [Apptainer](#) or [SingularityCE](#) — this is more common in High Performance Computing (HPC) academic clusters.

3.1.4.1 To run in `Docker` container

3.1.4.1.1 Using GitHub or Docker Hub container registries

- Clone the repository or download it manually and unpack.

```
git clone https://github.com/e-kotov/spanishoddata-paper-supplement.git
```

- To run the container image hosted in GitHub container registry run the following command while in the root directory of the current repository:

```
docker run --platform linux/amd64 --rm -p 8888:8888 -v $(pwd):/home/rstudio
ghcr.io/e-kotov/spanishoddata-paper-supplement:4.4.2
```

If the container is not available anymore in GitHub container registry (e.g. because of change of hosting terms), you can try getting the same container image from Docker Hub:

```
docker run --platform linux/amd64 --rm -p 8888:8888 egorkotovdhub/spanishoddata-
paper-supplement:4.4.2
```

If both container registries do not have the requested image, try downloading it using the [Zenodo hosted container](#) section below.

In terminal, look for the link 'http://127.0.0.1:8888/lab?token=SOMEALPHANUMERICSTRING' and open it in your browser. A JupyterLab will open in your browser. From there, click the `RStudio` button.

When in `RStudio`, run the following line to disable `renv`:

```
renv::deactivate()
```

The R session will restart and you will be able to run any pipeline using the packages that are already preinstalled in the container, no package installation from online source is necessary. To run the pipelines, kindly see the [relevant section below](#).

5

When you are done, press `Ctrl+C` in the terminal to stop the container.

3.1.4.1.2 Using Zenodo hosted container

• Clone the repository or download it manually and unpack.

```
git clone https://github.com/e-kotov/spanishoddata-paper-supplement.git
```

Download the container image file:

```
curl    https://zenodo.org/records/15207375/files/docker-container-image-r442-
pkg.tar.gz?download=1 --output docker-container-image-r442-pkg.tar.gz
```

Or download manually from https://doi.org/10.5281/zenodo.15207374.

Unpack the file:

```
gunzip -c docker-container-image-r442-pkg.tar.gz > docker-container-image-r442-
pkg.tar
```

Load the container image into Docker:

```
docker load -i docker-container-image-r442-pkg.tar
```

The image will be loaded with `r442spod` name, you can check that it was imported successfully with:

```
docker images
```

You can now run the container with:

```
docker run --platform linux/amd64 --rm -p 8888:8888 -v $(pwd):/home/rstudio
r442spod
```

In terminal, look for the link 'http://127.0.0.1:8888/lab?token=SOMEALPHANUMERICSTRING' and open it in your browser. A JupyterLab will open in your browser. From there, click the `RStudio` button.

When in `RStudio`, run the following line to disable `renv`:

```
renv::deactivate()
```

The R session will restart and you will be able to run any pipeline using the packages that are already preinstalled in the container, no package installation from online source is necessary. To run the pipelines, kindly see the relevant section below.

When you are done, press `Ctrl+C` in the terminal to stop the container.

3.2 To run in `Apptainer` container:

You can download the `Apptainer` container image from Zenodo https://doi.org/10.5281/zenodo.15207374.

The command to download would be:

```
curl https://zenodo.org/records/15207375/files/apptainer-container-image-r442-
pkg.sif?download=1 --output docker-container-image-r442-pkg.sif
```

You can then run the container with:

```
apptainer exec --bind "$(pwd)":/home/rstudio apptainer-container-image-r442-
pkg.sif bash
```

You may have to use extra commands and options to run Appatiner in High Performance Computing (HPC) academic clusters using SLURM job manager and bind the project folder, as well as forward the console output to a log file so that you can find the connection link. Kindly refer to the documentation provided by your HPC administrator as well as [Apptainer documentation](). Some HPC admins may offer a web-interface to run Apptainer containers using Jupyter Hub.

In terminal, look for the link 'http://127.0.0.1:8888/lab?token=SOMEALPHANUMERICSTRING' and open it in your browser. A JupyterLab will open in your browser. From there, click the `RStudio` button.

When in `RStudio`, run the following line to disable `renv`:

```
renv::deactivate()
```

The R session will restart and you will be able to run any pipeline using the packages that are already preinstalled in the container, no package installation from online source is necessary. To run the pipelines, kindly see the [relevant section below]().

When you are done, press `Ctrl+C` in the terminal to stop the container.

You can also use the same container to run on HPC as a developer container that you can connect to with SSH, as it has a built in SSH server, but you will need to configure it at start up time with a SLURM script and then tunnel/forward ports so that you can connect to it from Positron, VScode or Zed using their remote features.

4 Pipelines: original project steps

We provide all the code to reproduce the results in the paper, however we break it down into 4 steps. Step 1 caches some data on the date of the analysis, step 2 runs the speed test comparing CSV, DuckDB and parquet files, step 3 prepares the data for the plots, and step 4 generates the plots. Step 1 is not reproducible on a different date, as the data updates every day. Step 2 requires downloading a very large dataset and takes more then 24 hours to run. Step 3 is also rather resource intensive and downloads about 6 GB of data from the Internet.

All steps assume you have either local R or R running in the the provided containerised enviroment started in the project root directory.

4.1 1. Download and cache OSM data for Valencia

The line below will run the workflow to download the OSM data for Valencia and save it to the `data/proc/osm` directory.

```
Sys.setenv(TAR_PROJECT = "cache_osm_data"); targets::tar_make()
```

Because this data is changing every day, the downloaded file is then uploaded to https://doi.org/10.5281/zenodo.15207222, so that the data snapshot could always be found there.
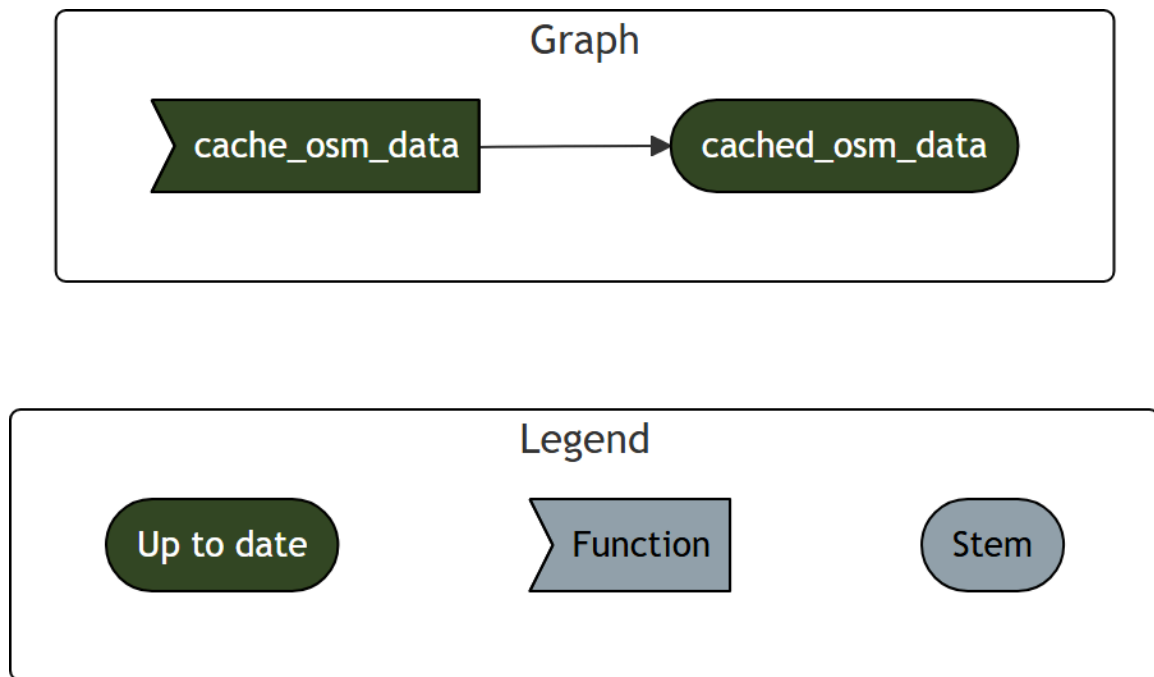

Figure 1: Pipeline visualisation of the `cache_osm_data` workflow

4.2 2. Run the speed test

The line below runs the workflow that downloads 18 GB of data using `spanishoddata` R package, converts it to `duckdb` and `parquet` (which will additionally take about 40 GB) and runs a series of tests with different number of threads and memory limits. This may take over a day and the combinations of threads and memory may need to be adjusted in the `_targets_speed_test.R` file in the root of the repository depending on your available hardware. Currently the maximum hardware specs to test are set up at 128 GB of memory and 64 processor cores.

```
Sys.setenv(TAR_PROJECT = "speed_test"); targets::tar_make()
```

The results are saved into `supplement-data/speed_test_results.csv` and are reused in step 3 below.

Figure 2: Pipeline visualisation of the `speed_test` workflow

4.3 3. Prepare the data for figures
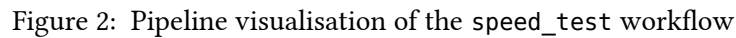
This step prepares all data, except the speed test data, to produce the article and supplement figures.

```r
Sys.setenv(TAR_PROJECT = "data_prep"); targets::tar_make()
```

Results are cached in the `_targets/data_prep` directory and are available in the repository, so that the final "main" pipeline can reuse them to generate the figures.

Figure 3: Pipeline visualisation of the `data_prep` workflow

4.4 4. Run the analysis and reproduce all figures

To run the analysis and reproduce all the figures run the following line:

```
Sys.setenv(TAR_PROJECT = "main"); targets::tar_make()
```

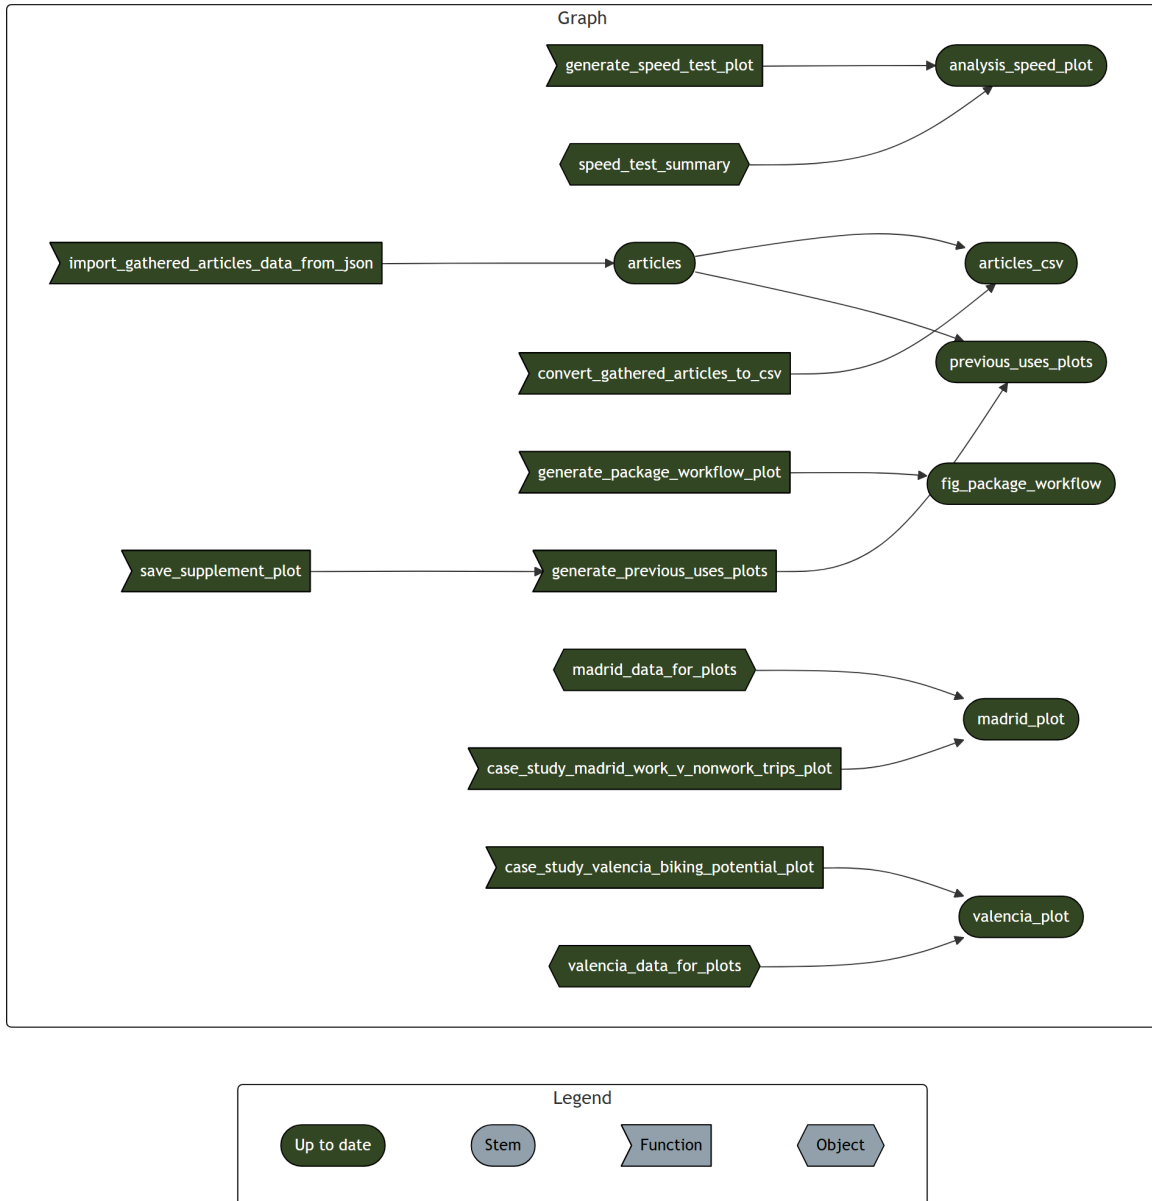The figures are generated in the `plots` directory.

Figure 4: Pipeline visualisation of the `main` workflow

4.5 Generate pipeline visualisations

To generate pipeline visualisations run the following line:

```
Sys.setenv(TAR_PROJECT = "pipeline-plots"); targets::tar_make()
```

The visualisations are generated as code for `mermaid` diagrams, as `html` and `png` files in the `media` directory. These are the plots you can observe above for each pipeline.

5 Building the computational environment from scracth

Here we provide some details how to build the environment from scratch.

5.1 Building the `Docker` container

Make sure `Docker` is installed and running. Open a terminal application in the project root directory.

```
docker build --platform linux/amd64 -f containers/Dockerfile -t r442spod .
```

To test run the container image locally, run the following command:

```
docker run --platform linux/amd64 --rm -p 8888:8888 -v $(pwd):/home/rstudio
r442spod
```

In terminal, look for the link 'http://127.0.0.1:8888/lab?token=SOMEALPHANUMERICSTRING' and open it in your browser. A JupyterLab will open in your browser. From there, click the `RStudio` button.

When you are done, press `Ctrl+C` in the terminal to stop the container.

To archive the container image to a file:

```
docker save r442spod -o ~/docker-container-image-r442-pkg.tar
```

5.2 Building the `Apptainer` container

The resulting `tar` file with the `Docker` container image can be converted to an `Apptainer`/ `Singularity` container.

If you are on HPC, make sure to enable the `apptainer` module. This might be done with:

```
module load apptainer
```

Assuming the Docker image is saved as `docker-container-image-r442-pkg.tar` and uploaded to ~/ on HPC:

```
apptainer build ~/apptainer-container-image-r442-pkg.sif ~/docker-
archive:docker-container-image-r442-pkg.tar
```

To run the `Apptainer` container in HPC you might have to set up a `SLURM` job. Kindly, follow the instructions provided by your HPC adminstrator.